

PROJECTS

- **Pickup Tab Experience in Amazon Detail Page**
 - Implemented a "*Collection Tab*" in the Product Page where Pickup Options are shown to customers if eligible.
 - Created an Event-driven, Pre-Compute system in AWS (CDK-based) that maintains all reqd. Pickup Location Metadata. This system processes 2K+ CRUD events per second.
 - Implemented the delivery date messaging (*FREE pickup Tuesday, July 2*) for the **Pickup Tab** in the Amazon Product Page.
 - It was a complex problem which could have introduced additional latency to the entire Product Page.
 - Worked with Principal Engineers across Retail org. and arrived at a solution which did NOT add any Latency/Traffic overhead.
 - Worked with other core-Amazon teams to integrate with their stack to compute and render necessary data for Pickup Offers(Price, Item Availability, Merchant Details etc.).
 - The *Tab* is rendered upto 40K times in the core Product Page and uplifted orders to Pickup Locations by an estimated 10% (~8K orders per day on average).
- **DHL Postnumber Validation**
 - Amazon has a tie up with DHL in Germany where customers can order to DHL Pickup Locations from [amazon.de](#).
 - To order to DHL pickup locations, customers required a mandatory id called the "*postnumber*".
 - However, customers sometimes entered incorrect/garbage values for their *postnumber* (ex - 123). This was leading to almost ~5K failed deliveries everyday across Germany.
 - Reached out and collaborated with DHL's engineers to implement DHL's internal validation algorithm.
 - Launched the validations which now entirely prevents incorrect postnumbers from being used.
- **Accurate Estimated Delivery Date Messaging for Pickup Locations in Poland**
 - There was a *long existing gap* in the Delivery date service that caused a 1-2 day delay in the estimated delivery date on the Product Page. If the order could arrive by *Tuesday*, the estimated date would be on a *Thursday*.
 - A good portion (~20%) of [amazon.pl](#) 's overall orders are delivered to Pickup Locations.
 - This was marked as a critical blocker to launch *Amazon Prime* in amazon.pl (due to high usage of Pickup Locations and Amazon Prime's requirement to provide faster shipping).
 - Identified the gap and spearheaded the development of the solution across *5 core Amazon teams*. The primary challenge was to ensure the solution met the diverse requirements and constraints of all teams, including service tenets, latency/TPS, and operational overhead. Successfully aligned all teams to reach a consensus on a solution that satisfied these constraints.
 - Worked with the amazon.pl Vice President(s) to escalate and prioritize the progress for the fix. Delivered the fix before the expected date and enabled the launch of Prime in Poland.
- **Pickup Checkout Experience WW**
 - Implemented a low-latency (<10 ms) API used in Checkout to "upsell" customers about Amazon Hub Locations near their addresses. This API serves upto ~7K customer requests per second.
 - Owned the primary discoverability microservice across design reviews, office hours, New Feature Requests, etc.
 - Made multiple optimizations and improvements to the microservice -
 - Identified a redundant call pattern reducing ~300 TPS in downstreams traffic for Pickup Location Eligibility.
 - Built Canaries ("OneBoxes") to catch potential issues early, during deployments.
- **Operational Excellence**
 - Enabled dynamic distributed throttling to eliminate bot-related service brownouts. Implemented comprehensive testing and configured dashboards and alarms. Successfully managed peak readiness for Black Friday and Prime Day from 2022-2024, ensuring optimal

scaling and performance while coordinating traffic expectations with upstream and downstream teams.

- **Pickup Config Centralization and Onboarding Automation**

- There are multiple teams under the Pickup Charter with distinct responsibilities like - Metadata, Discoverability, Ranking, Post-Order, Notifications, Eligibility/Capacity Checks, etc.
- Multiple teams were using their own data stores to maintain *Partner specific configuration*. As a result, if a configuration had to be added/deleted, it was a cross-team manual effort that took up to 4 developer days.
- Worked with all involved teams to identify the Configuration requirements and came up with an approach to centralize the configuration to a single data-store. Created a common access library for the centralized data-store which allows each team to read the configurations.
- Since the configuration was critical to the proper functioning of multiple services/teams, gated the UPDATE / DELETE operations via a *new microservice*. Added human-specific and service-specific authorization and authentication, enabling teams to modify their respective data.
- Setup a call-chain across all Microservices to automate the configuration onboarding process
- With this solution, the Partner was only required to be onboarded at a single place.

- **Return to Pickup Point Optimization**

- This was a pilot program whose initial design had a huge scaling bottleneck.
- With every expansion to additional zip codes, we required scaling across 3 different microservices and 2 different teams.
- It was a tedious process as traffic varied from zip code to zipcode based on customer behavior (which couldn't be predicted). We were always either over-scaled or under-scaled with each expansion and had to correct our scaling with time, costing a lot of developer effort and infrastructure costs.
- Proposed a modification to the existing design by introducing a tradeoff that slightly reduced accuracy for a significant reduction in the overall load to our downstream services.
- Derived metrics from the existing customer experience to quantify the potential impact (3% drop in Accuracy). Used this to back up the proposal and got the alignment of all Product Managers involved. Implemented the same.
- Since then, scaling was only needed 25% of the time when compared to earlier and was also only limited to a single microservice.